

Hidden Markov Music

Daniel Wysocki

QUEST, April 15, 2015



Knowledge-based Systems

- follow a set of rules defined by the programmer
- depends on knowledge of the programmer



Machine Learning

- existing compositions are used to create a model
- new compositions are produced based on the model
 - deterministic
 - probabilistic
- challenging to find a model which captures the essence of music



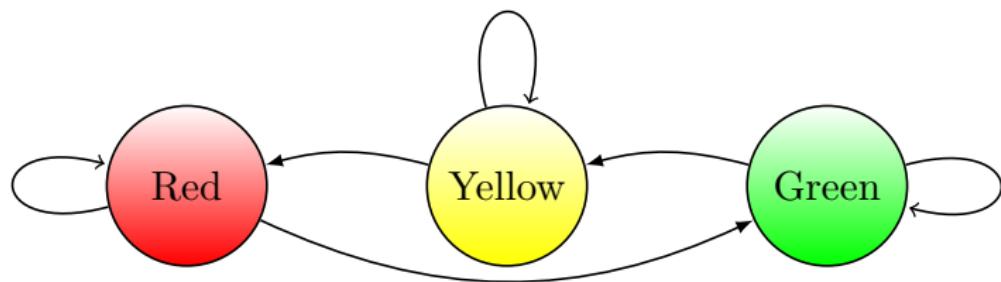
Markov Processes



Definition

- the future depends only on the present
- nondeterministic
- may not perfectly represent the system being modeled
 - often serves as a good approximation

Markov Chain



Training a Markov Chain

- to train a Markov chain, simply count the occurrences of each transition
- divide each element by its row's total

	G	Y	R
G	45	5	0
Y	0	25	25
R	30	0	20

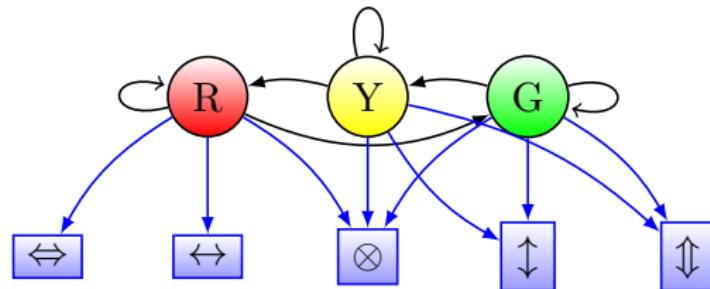
⇒

	G	Y	R
G	0.9	0.1	0
Y	0	0.5	0.5
R	0.6	0	0.4

Hidden Markov Model

- Marvin the Martian is looking down at a traffic light from space
- he cannot see the actual lights, but instead he sees the speed and direction of the cars
- he can still model the traffic light, using an HMM

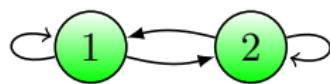
Hidden Markov Model Example



Overview

- we model songs as Markov processes
- notes are observed
- some underlying states of the song are hidden from us
 - we choose the number of states, and everything else is automatic
- we train the model on a song
 - allows us to generate new songs (algorithmic composition)

Model



C

D

E

F

G

A

B

END

First Song

- trained a model on Twinkle, Twinkle, Little Star [Play](#)
- produced the following song [Play](#)



Für Elise

- trained a model on Beethoven's Für Elise [Play](#)
- using 5 states [Play](#)
- using 15 states [Play](#)

Future Improvements

- incorporate duration into the notes in an intelligent way
- combine songs from multiple composers into a single model



Acknowledgements

Special thanks to Craig Graci, and Andrey Markov



Questions?

Listen to more here



<https://dwysocki.github.io/csc466/music.html>

